

Graph Partitioning

Introduction


1. Common thread in a distributed memory algorithm: How do you distribute the data?
 - Spectral partitioning: Exploits the connection between graphs and linear algebra
 - Physics-based interpretation based on systems of springs

The Graph Partitioning Problem

1. Can represent a graph as an adjacency matrix
 - Can partition the matrix across rows, corresponding to a graph (vertex) partition (sparse matrix)
 - Can then do a BFS with a matrix-vector multiply
 - Implies partitioning the vector as well
 - Amount of work corresponds to the number of non-zero entries
 - Ideally, we should balance the work
 - Any time an edge crosses a process boundary, a communication exchange occurs
 - Reduce edge cuts to minimize communication volume
2. Two goals:
 - Balance work across processes
 - Minimize communication between processes
3. Graph partitioning problem:
 - Given $G = (V, E)$ and number of partitions, P
 - Compute a (vertex) partition $V = V_0 \cup V_1 \cup \dots \cup V_{p-1}$ such that:
 - $\{V_i\}$ are disjoint $\Rightarrow \text{intersect}(V_i, V_j) = 0$
 - $\{V_i\}$ are roughly balanced $\Rightarrow |V_i| \sim |V_j|$
 - Let $E_{\text{cut}} = \{(u, v) \mid u \text{ in } V_i, v \text{ in } V_j, i \neq j\}$ and minimize E_{cut}

Do You Really Want a Graph Partition?

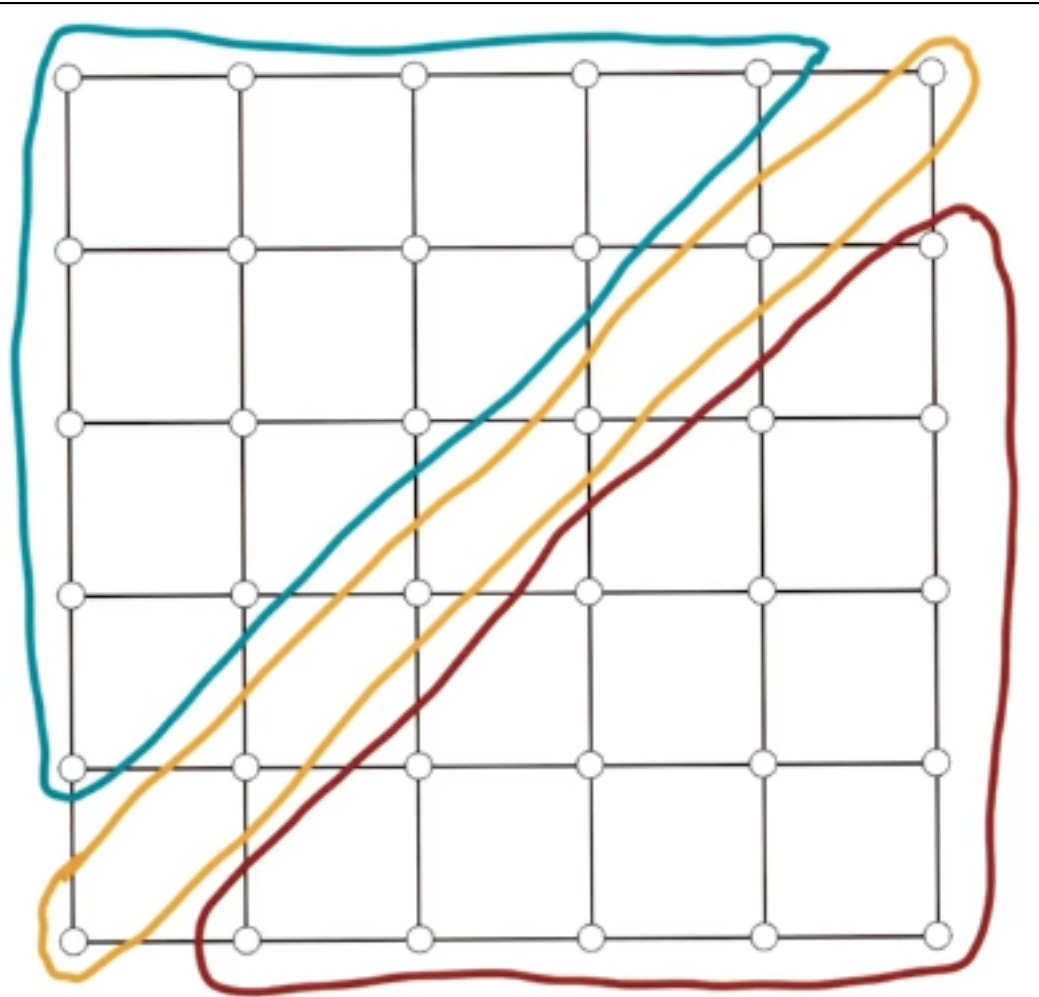
1. Does the graph partitioning problem as formalized above meet the intended goals?
 - No. Despite the number of vertices per partition being the same, the number of non-zero entries can vary, so the work is not evenly distributed



Uneven Distribution of Work

Graph Bisection and Planar Separators

1. Graph partitioning is NP-complete -> Need heuristics!
 - Heuristic: Bisection (divide and conquer)
 - Suppose you want P partitions
 - Start by using any algorithm to divide the graph into 2 partitions
 - Then, divide each half into two more partitions
 - Repeat until the P partitions are obtained
2. Planar graphs
 - Graphs that can be drawn in 2D with no edge crossings
3. Planar separators
 - Theorem: A planar graph $G = (V, E)$ with $|V| = n$ vertices has a disjoint partition $V = A \cup S \cup B$ such that:
 - S separates A and B
 - $|A|, |B| \leq 2n / 3$ (the larger of the two partitions is no more than twice the size of the smaller)
 - $|S| \leq 2 * \sqrt{2n} = O(\sqrt{N})$
 - Lipton and Tarjan 1979
 - Lipton and Tarjan described a polynomial-time algorithm



Planar Graph

Partitioning via Breadth First Search

1. How can BFS be used to bisect a graph? Assume the graph is connected
 - Run BFS from any vertex
 - BFS will assign every vertex to a level
 - Levels separate subgraphs
 - Stop when about 1/2 of the vertices have been visited. Assign visited to one partition, unvisited to the other
 - Other criteria are possible
2. However, we wanted to use graph partitioning to distribute our BFS computation, so we can't use BFS to determine the graph partition

Kernighan Lin - Part 1: No Gain is Pain

1. Kernighan-Lin algorithm is the most well-known heuristic for graph partitioning
 - Given a graph, divide the vertices into two subsets of equal or nearly- equal size
 - $V = V1 \cup V2$, $|V1| = |V2|$
 - Cost of this partition is the number of edges between $V1$ and $V2$

- $\text{Cost}(V1, V2) = \text{Number of edges between } V1 \text{ and } V2$
 - Assume you have two evenly-sized subsets of $V1$ and $V2$
 - $X1$ from $V1$, $X2$ from $V2$ with $|X1| = |X2|$
 - If you swapped these two subsets, the cost will change, but by how much?
2. Formal definitions
- Consider vertex a in $V1$ and vertex b in $V2$
 - External costs: Number of vertices in the opposite partition
 - $E1(a \text{ in } V1) = \# \text{ edges}(a, b \text{ in } V2)$
 - $E2(b \text{ in } V2) = \# \text{ edges}(b, a \text{ in } V1)$
 - Internal costs: Number of vertices in the same partition
 - $I1(a \text{ in } V1) = \# \text{ edges}(a, i \text{ in } V1)$
 - $I2(b \text{ in } V2) = \# \text{ edges}(b, i \text{ in } V2)$
 - $\text{Cost}(V1, V2) = \text{Cost}(V1 - \{a\}, V2 - \{b\}) + E1(a) + E2(b) - c(a,b)$
 - c is a constant to account for an edge between a and b . 1 if an edge exists, 0 otherwise
 - Then, consider the cost after swapping a and b
 - $\text{Cost}(V1, V2) = \text{Cost}(V1 - \{a\}, V2 - \{b\}) + I1(a) + I2(b) + c(a,b)$
 - Change in cost = $E1(a) + E2(b) - I1(a) - I2(b) - 2c(a,b)$
 - Larger change is better \Rightarrow Larger decrease in cost
 - Change in cost = $\text{gain}(a \text{ in } V1, b \text{ in } V2)$
 - Can be negative if cost increased

Kernighan Lin Algorithm Quiz

1. Consider the computation of the gain function a and b
 - Assume the following:
 - Every vertex has a partition label; $O(1)$ access time
 - Maximum degree of any vertex is d
2. What is the sequential running time to compute $\text{gain}(a, b)$ in terms of d , $n1 = |V1|$, $n2 = |V2|$
 - $O(d)$
 - Need to sweep over the adjacent vertices, which is at most d
 - Need to check partition label of each vertex, which is a constant-time operation

Kernighan-Lin Algorithm

1. How do we choose $X1$ and $X2$, the subsets of the partitions?

```

let V = V1 U V2
C = cost(V1, V2)
forall a in V1, do compute E1(a), I1(a)
forall b in V2, do compute E2(b), I2(b)
forall v in V do visited[v] <- false

while !all(visited) do
    choose unmark (a in V1, b in V2) with largest gain(a, b)
    visited[a], visited[b] <- true
    Update all E1, E2, I1, I2 // not actually swapping nodes, just updating costs

// this provides gain(a1, b1), gain(a2, b2), ...
let Gain(j) = sum(gain(a, b))
choose jmax = argmax(Gain(j))
if Gain(jmax) > 0 then
    X1 = {a1, a2, ... ajmax}
    X2 = {b1, b2, ... bjmax}
    Update C <- C - Gain(jmax)
    V1 <- (V1 - X1) U X2 // swap


```

```
V2 <- (V2 - X2) U X1 // swap
```

```
// repeat until Gain(jmax) < 0
```

2. Cost:

- Overall running time $O(|V|^2 * d)$
- Can be reduced to $O(|E|)$
 - Fiduccia and Mattheyses 1982



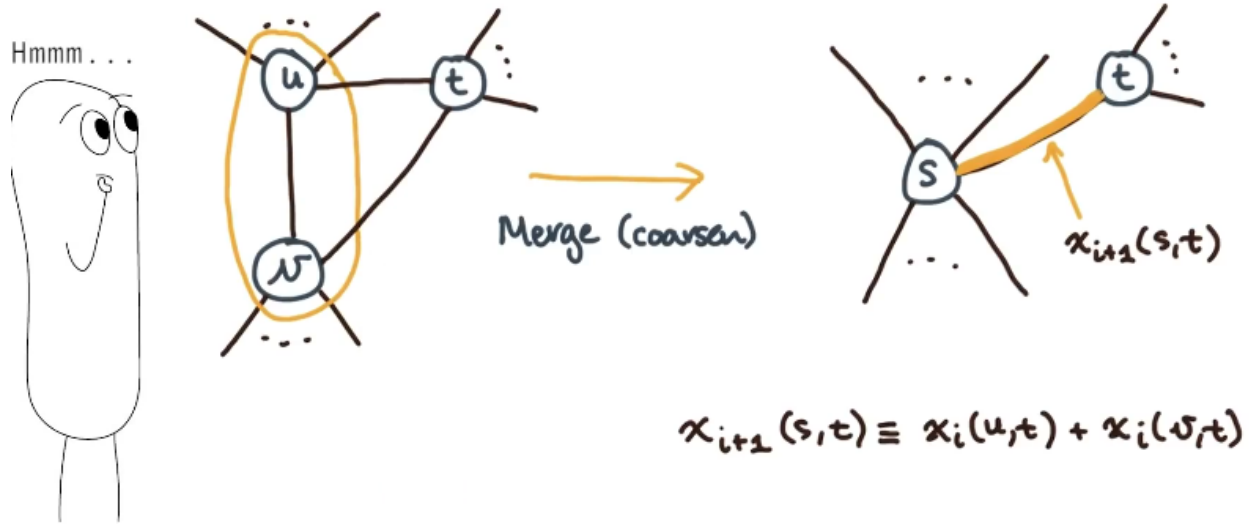
Kernighan-Lin Algorithm

Graph Coarsening

1. Multi-level graph coarsening: Different partitioning scheme
 - Divide-and-conquer
 - Continually replace graph with smaller graphs (fewer nodes and edges) but still somehow resembles original graph
 - Repeat until the graph is small enough to partition quickly
 - If we've done a good job of preserving the shape of the graph, the split will correspond to a roughly equivalent split in the larger graph
2. How do you actually coarsen a graph?
 - Identify at least one subset of the vertices to collapse
 - Replace subset with single "super" vertex
 - Track the fact that one vertex has replaced many by assigning a weight W to the vertex, where W is the number of replaced vertices
 - Track edge weights as well so we can cut edges accurately later on
 - $G_i = (V_i, w_i: V_i \rightarrow R, E_i, x_i: W_i \rightarrow R)$
 - w_i and x_i are functions that map edges and vertices to super edges and vertices
 - Simply sum the contributions from each vertex

$$G_i = (V_i, w_i: V_i \rightarrow \mathbb{R}, E_i, x_i: E_i \rightarrow \mathbb{R})$$

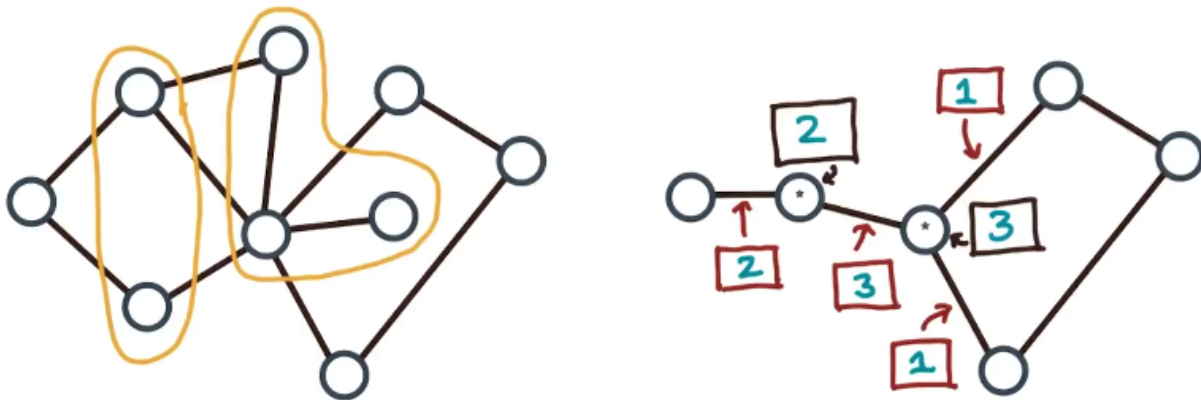
$$G_{i+1}$$



Graph Coarsening Algorithm

Coarsen Me - Baby

1. Determine the coarsen vertex and edge weights in the following graph:



Graph Coarsening Quiz

Maximal and Maximum Matchings

1. Matching: A subset of a graph $G = (V, E)$ is a subset E' of E of edges with no common endpoints
 - Independent set, but for edges instead of vertices
2. Maximal matching: A matching is maximal if no more edges may be added
3. Maximum matching: A graph's maximum matching is its largest (most edges or total edge weight)

Def.: A graph's **maximum** matching is its largest (most edges or total edge weight).

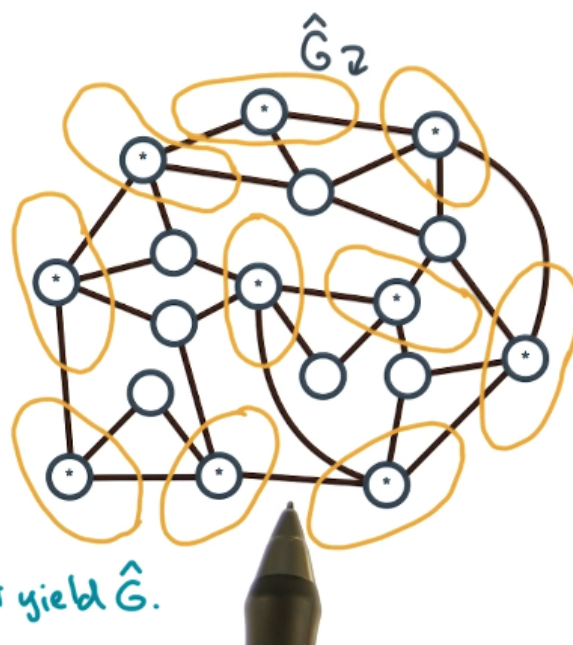
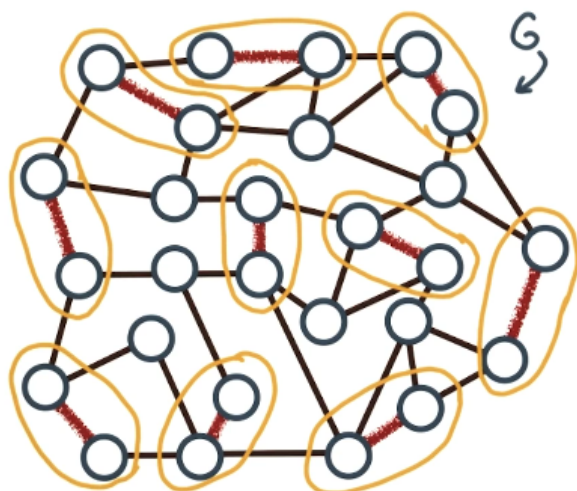


Maximal (and Maximum) Matching

Find the Maximal Matching

1. Mark the edges of G that yield G_{hat}

Quiz! Find the Maximal Matching



Your task: Mark the edges of G that yield \hat{G} .

Maximal Matching Quiz

A Fact about Maximal Matchings

1. Consider an initial graph G . Suppose you find a maximal matching and use it to coarsen. Repeat this process producing a sequence of graphs from 1 to k . The original graph has n vertices and the final graph has s vertices.
2. How large must k be in terms of n and s ?
 - $\log(n/s)$
 - Each level must have at least half of the previous number of vertices, so $|V_k| = n/2^k \Rightarrow k \geq \log_2(n/s)$


Computing a Maximal Matching

1. Pick an unmatched vertex at random
2. Match to one of its unmatched neighbors
 - Could pick randomly
 - Better strategy is to pick the heaviest edge
 - Not a lot of rigorous analysis, but lots of experimental evidence
 - The intuition comes from the fact that picking the heaviest edge should lead to the greatest reduction in edge weight when coarsening the graph

Fine-to-Coarse and Back Again

1. Mark the planar separator from the coarsened graph in the original graph
 - Some separator edges map ambiguously if they were merged in the coarsened graph

Quiz! Fine-to-Coarse & Back Again



Your task: Mark the separator in the original (left) graph.

Planar Separator in Coarsened Graph

Partition Refinement

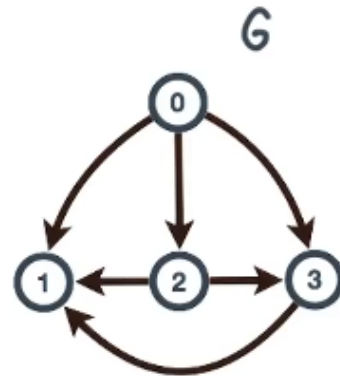
1. A minimum balanced edge cut in a coarsened graph minimizes the balanced edge cut in the next finer graph.
 - False; Coarsening is a heuristic, so it's possible that there's a better cut in the finer graph

Spectral Partitioning - Part 1: The Graph Laplacian

1. Instead of representing a graph as an adjacency matrix, represent it as an incidence matrix
 - Each row is an edge, each column is a vertex
 - Graph Laplacian, $L(G) = C' * C$
 - Diagonals: Count incident edges (always 1)
 - Off-diagonals: Says edges are adjacent
 - Tells us something about the undirected form of the original graph; we lose the direction information when calculating the Laplacian
 - Graph Laplacian, $L(G) = C' * C = D - W$
 - D is the degree of each vertex (along the diagonal)
 - W is the adjacency matrix marking all the edges

$C = C(G)$: Incidence matrix

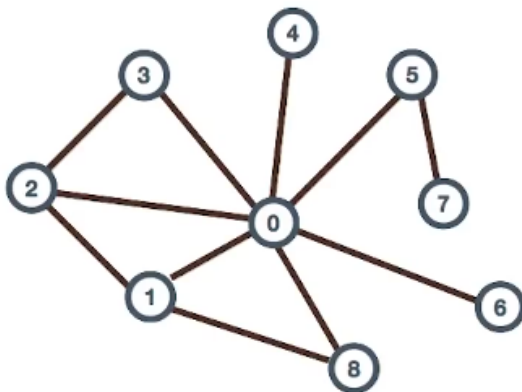
	0	1	2	3	
0	1	-1			$0 \rightarrow 1$
1	1		-1		$0 \rightarrow 2$
2	1			-1	$0 \rightarrow 3$
3		-1	1		$2 \rightarrow 1$
4			1	-1	$2 \rightarrow 3$
5		-1		1	$3 \rightarrow 1$



Incidence Matrix Representation

Graph Laplacian

1. Compute the graph Laplacian for the following graph
 - Matrix should be symmetric due to the definition of the Laplacian



	0	1	2	3	4	5	6	7	8
0	7	-1	-1	-1	-1	-1	-1		-1
1	-1	3	-1						-1
2	-1	-1	3	-1					
3	-1		-1	2					
4	-1				1				
5	-1					2	-1		
6	-1						1		
7						-1	1		
8	-1	-1							2


Laplacian Quiz 1

Spectral Partitioning - Part 2: Springs Fling

1. If we imagine a series of weights connected with springs displaced from a common point x , it provides some physical intuition for the Laplacian
 - Force on one weight is proportional to the displacement of the adjacent weights
 - This corresponds to the graph Laplacian for a line graph

A 2D Laplacian

1. What is the graph Laplacian for a group of 9 weights connected in the following way:

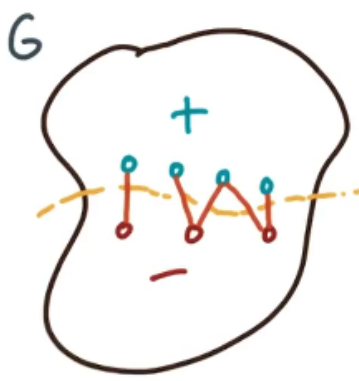


Laplacian Quiz 2

Spectral Partitioning - Part 3: Algebraic Connectivity

1. Factoids

- Laplacian $L(G)$ is symmetric
- $L(G)$ has real-valued, non-negative eigenvalues and real-valued, orthogonal eigenvectors (not complex)
 - Eigenvalues and eigenvectors are pairs; multiplying $L(G)$ by its eigenvector gives a scaled version of the eigenvector. The scaling factor is the eigenvalue
 - $L(G) * Q = Q * Y$ where Q and Y are matrices
 - Columns of Q are the eigenvectors and diagonal entries of Y are the eigenvalues
 - Convention: Assume we can sort the eigenvalues
- G has k connected components if and only if the k smallest eigenvalues are identically 0
 - Spectrum of $L(G)$ tells us something about the connectivity of G
- Let $V+$ be the number of vertices in the positive section and $V-$ be the number of vertices in the negative section
 - Define x such that x_i is 1 if i is in $V+$ and -1 if i is in $V-$
 - Then, the number of cut edges = $1/4 * x' * L(G) * x$

G


let $V_+ \equiv \{\text{vertices in "+"}\}$
 $V_- \equiv \{\text{vertices in "-"}\}$


let $\vec{x} \equiv \begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_{n-1} \end{bmatrix}$ s.t. $x_i = \begin{cases} +1 & \text{if } i \in V_+ \\ -1 & \text{if } i \in V_- \end{cases}$

Fact 4) # of **cut edges** = $\frac{1}{4} \vec{x}^T L(G) \vec{x}$

Spectral Partitioning

Counting Edge Cuts

Quiz! Counting Edge Cuts

G


Your task:
 Fill in the boxes!

$$\begin{aligned}
 \vec{x}^T L(G) \vec{x} &= \sum_{i,j} l_{ij} x_i x_j \\
 &= \sum_{i,j} l_{ij} x_i x_j \quad \leftarrow = l_{ii} x_i^2 \quad \boxed{d_i} \quad \boxed{+1} \\
 &+ \sum_{\substack{i,j \in V_+ \\ i \neq j}} l_{ij} x_i x_j \quad \boxed{-1} \quad \boxed{1} \quad \boxed{1} \quad + \sum_{\substack{i,j \in V_- \\ i \neq j}} l_{ij} x_i x_j \quad \boxed{-1} \quad \boxed{1} \quad \boxed{1} \\
 &+ \sum_{\substack{i \in V_+ \\ j \in V_-}} l_{ij} x_i x_j \quad \boxed{-1} \quad \boxed{1} \quad \boxed{-1} \quad + \sum_{\substack{i \in V_- \\ j \in V_+}} l_{ij} x_i x_j \quad \boxed{-1} \quad \boxed{-1} \quad \boxed{1}
 \end{aligned}$$

Laplacian Quiz 3

Spectral Partitioning - Part 4

1. Start with a graph $G = (V, E)$
2. Construct its Laplacian $L(G) = D - W$
3. Suppose we have a partitioning of the vertices $V = V_+ \cup V_-$
4. Translate this into a partition vector $\vec{x} = \{+1 \text{ if } i \in V_+, -1 \text{ if } i \in V_-\}$

5. If we want to minimize the cut edges, we need to pick an x to minimize $1/4 * x' * L(G) * x$
 - Want: $\min(1/4 * x' * L(G) * x)$
 - $\sum(x_i) = 0$ (same number of vertices in each partition)
 - $x_i = +1$ or -1
 - These constraints make the problem NP-complete
 - If we relax the constraint that we assign a $+1$ or -1 to every vertex, we can use the Courant-Fisher Minimax theorem
 - This says that the vector that minimizes this problem is q_1 , which is the eigenvector corresponding to the second smallest eigenvalue of $L(G)$
 - Choosing x to be q_1 gives us a lower bound, but how do we take q_1 and turn it into a partition vector?
 - Based on the spring model, the second smallest eigenvector has some sine-like shape where half the values are positive and half are negative
6. Spectral partitioning algorithm:
 - Create $L(G)$
 - Compute (Y, q_1) eigenpair of $L(G)$
 - Choose $x(i) = \text{sign}(q_1(i))$
7. Spectral partitioning works very well for planar graphs

Conclusion

1. Graph partitioning is NP-complete
 - Want algorithms with good heuristics
2. Heuristics:
 - Multilevel partitioning (divide and conquer)
 - Exploiting special structure (planarity)
 - Improvement techniques (Kernighan-Lin)
3. Spectral partitioning: Exploits relationship between graphs and matrices
 - K-L and spectral partitioning are very expensive relative to the motivation, which was a relatively cheap BFS