

# Introduction to High Performance Computing

## Welcome

1. Explore big ideas about how to extract parallelism and data locality from algorithms and data structures

## Readiness Survey

## Philosophy and Logistics

1. Lectures: Develop intuition
2. Readings: Formalize intuition mathematically
3. Projects: Writing real code
  - Don't understand an algorithm or data structure until actually implemented

## What is HPC?

1. High performance is an ambiguous term because it's relative
  - Supercomputing is also used to refer to the same concepts

## Supercomputing in a Nutshell

1. Given a computational problem and a machine, how do you compute at the absolute limits of scale and speed?
  - Limits come from physics (speed of light or amount of energy, power, or heat dissipated by a system)

## Topics

1. Sequential or serial RAM model: Single serial processor connected to memory
  - Issues instructions that operate on data
  - Operands always live in memory
  - Assume all instructions have a cost that is bounded by some constant
    - $O(n \log(n))$
2. Parallel RAM model: Multiple processors connected to shared memory
  - Still assume constant cost per operation
  - Processors can communicate via shared variables
  - Reduce total cost by up to  $P$ , where  $P$  is the number of processors
    - $O(n * \log(n) / P)$
3. Distributed memory network model: Interconnected network of RAMs
  - Each processor has separate memory, no processor can read or write to any other processor's memory
    - Instead, must communicate over a network
  - $\text{Cost}(n, P) = a * f(n, P) + B * g(n, p)$ 
    - Need to account for number of messages and total volume of communication
4. Two-level I/O Memory: Multiple processors share memory, but there's at least one additional level of faster memory sitting between the processor and the slower main memory
  - How much data needs to move between the memory and processor using the intermediate "scratch pad" memory
  - Cache, virtual memory