

Shared Memory Parallel BFS

Introduction

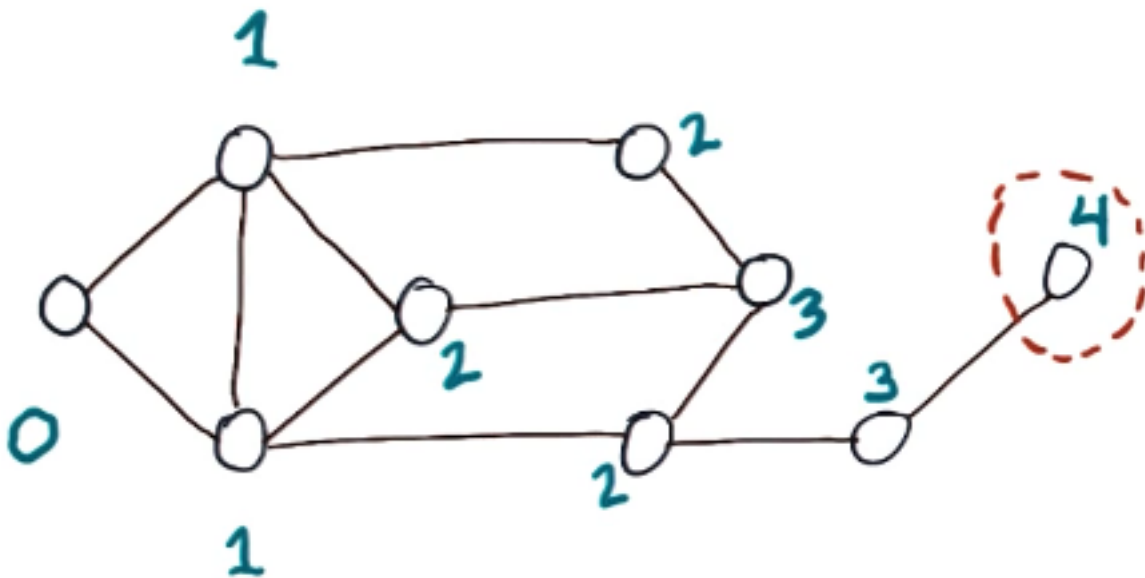
1. Graphs are essential in modern computing
 - Social networks, roads, power grids
2. This lesson examines a parallel breadth-first search algorithm for the dynamic multithreading model

BFS 101

1. Breadth-first search: Given a graph and a starting vertex, what is the distance to all other nodes in the graph?
2. Work and span
 - $W(n) = O(|V| + |E|)$
 - We visit each vertex at most once and each edge at most once

BFS($G=(V,E)$, s)

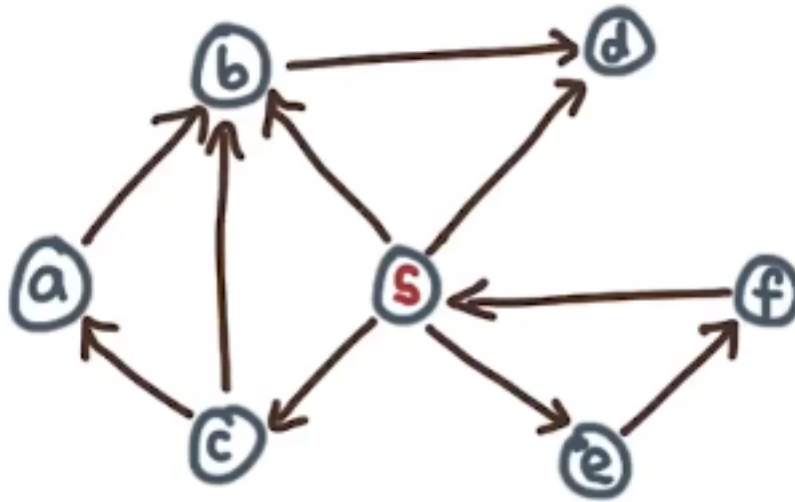
```
{
  D[V] <- inf
  D[s] <- 0
  F <- {s} // queue of unvisited vertices
  while F != 0 do
    v <- extract_one(F)
    for(v,w) <- E do
      if D[w] = inf then
        D[w] <- D[v] + 1
        F <- F U {w}
  return D // D[x] = dist(s -> x)
}
```



BFS Example

BFS Example

1. Consider the following directed graph:



BFS Quiz

2. What does F contain right after the third execution of the while loop? There are multiple correct answers due to the ambiguity in “extract_one”
 - c d e

Is BFS Inherently Sequential?

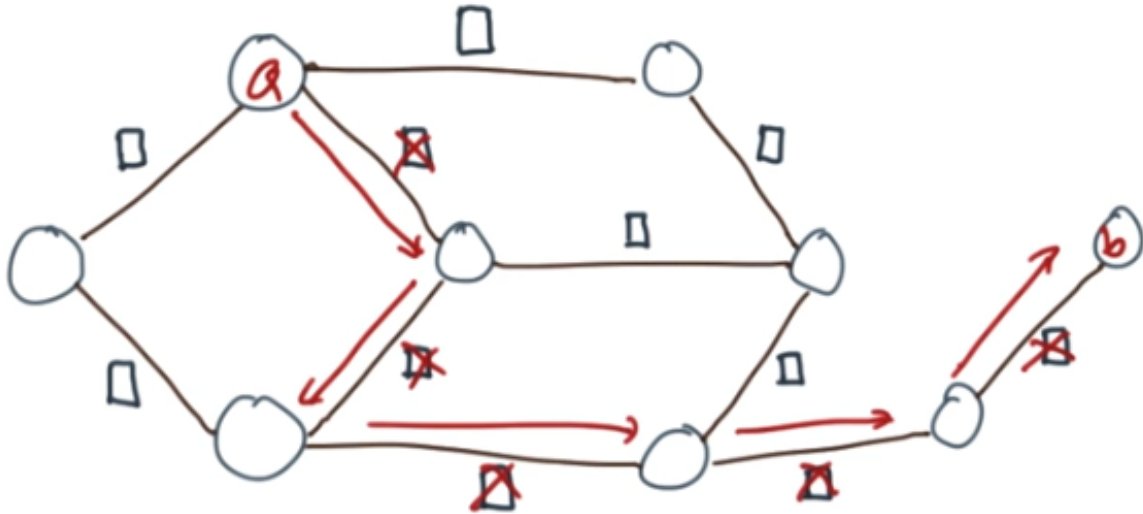
1. BFS Work and Span
 - $W(n) = O(|V| + |E|)$
 - $D(n) = O(|V|)$
 - $W/D = O(1 + |E|/|V|)$
 - In real life, graphs are sparse which means $|E| = O(|V|)$
 - This means the average available parallelism will be a constant, which is no good

Intuition - Why We Might Do Better

1. The upper bound on the span of BFS shouldn't be the number of vertices, it should be the number of waves
 - Waves are sets of nodes that are reachable from a different set of nodes
 - Diameter: Maximum shortest distance between any pair of vertices in a graph
2. Level-synchronous traversal: Visit nodes level by level
 - Doesn't matter which order we visit nodes within a level
 - This means we can look at nodes simultaneously -> parallelism

BFS Example 2

1. Find the path with 5 edges in the previous example



BFS Quiz

High Level Approach to Parallel BFS

1. Two key ideas:
 - Level synchronous
 - Process an entire level in parallel
2. processLevel takes the graph and current frontier and creates a new frontier, as well as updates the distances
 - Frontiers are level-specific waves of the graph
 - l is a level counter
 - Use a special data structure called a bag for the frontiers

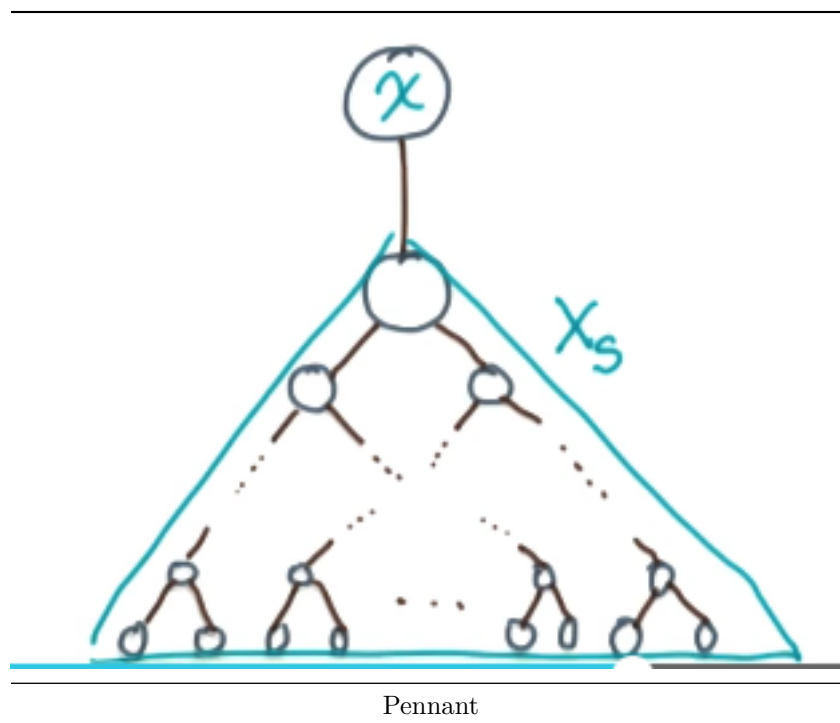
```
parallelBFS(G=(V,E), s)
{
  D[v] <- inf
  D[s] <- 0
  l <- 0
  F0 <- {s}
  while F1 != 0 do
    F(l+1) <- {}
    processLevel(G,F1,F(l+1),D)
    l <- l + 1
  return D
}
```

Bags: Key Properties

1. Container having the following properties:
 - Unordered collection
 - With repetition
2. Operations
 - Fast traversal
 - Fast and logically associative ($A \cup B == B \cup A$)
 - “Union” and “split” should allow us to divide the data structure into roughly equal pieces and combine them back together

Pennants - Building Blocks for Bags

1. Pennant: Tree with 2^k nodes and a unary root have a child that is a complete binary tree
 - Root: x
 - Child: Xs
2. Consider two same-sized pennants x and y with children Xs and Ys respectively
 - To combine, make x the new root, y the child of x , and Xs and Ys the children of y
 - Gives a pennant of size $2^{(k+1)}$
 - We can undo this operation just as easily by repeating the same steps in reverse



Pennants

1. What is the output of combining two pennants of different sizes?
 - This is invalid, pennants must be the same size to combine

Combining Pennants Into Bags

1. Represent 23 in binary = 10111
 - You can use this to split a bag into pennants by creating pennants of size 1, 2, 4, and 16 (in the case of 23)
 - Make an array of pointers to the start of each pennant (called a “spine”)

Duality Between Bags and Binary Math

1. $23 + 1 = 10111 + 1 = 11000$
2. If we want to add an element to a bag, we follow the same approach as base-2 addition
 - Carry pennants that are the wrong size into the next slot of the spine

What is the Cost of Insertion?

1. What is the cost of inserting one element into a bag?

- $O(1)$
 - $O(\log(n))$ (this one)
 - $O(n)$
 - $O(n * \log(n))$
2. In the worst case, we must traverse the entire spine to insert
 - The spine requires $\text{ceil}(\log_2(n))$ elements to represent an integer n , so the complexity is $O(\log(n))$
 - Each insertion requires constant time to perform

Combining Two Bags

1. What is the cost of bag union?
 - $O(1)$
 - $O(\log(n))$ (this one)
 - $O(n)$
 - $O(n * \log(n))$
 - $O(n^2)$
2. The same logic applies as above
 - This is cool because it takes the same amount of time to insert n elements into a bag as it does to insert one element
 - The amortized cost is a constant

Cost of Other Operations

1. What is the cost of bag splitting?
 - $O(1)$
 - $O(\log(n))$ (this one)
 - $O(n)$
 - $O(n * \log(n))$

Bag Splitting

1. Performing a right bit shift is roughly equivalent to dividing by two
 - $23 = 10111 \gg 1 = 1011 = 11$
2. Cost = $[O(\log(n) \text{ splits}) * [O(1) \text{ per split}] = O(\log(n))$

Finishing the Parallel BFS with Bags

1. The following pseudocode implements the processLevel function for the parallel BFS algorithm stated above:

```
processLevel(G, Fl, Fl+1, D)
{
    if |Fl| > phi then
        (A,B) <- bagSplit(Fl)
        spawn processLevel(G, A, Fl+1, D)
        processLevel(G, B, Fl+1, D)
        sync
    else
        for v in Fl do
            parfor (v,w) in E do
                if D[w] = inf then
                    D[w] <- 1 + 1
                    bagInsert(Fl+1, w)
}
```

2. Even though there's a data race in the base case where multiple threads might write to the same $D[w]$ location, because the algorithm is level-synchronous, each thread is guaranteed to write the same value
3. Work
 - $W(n) = O(|V| + |E|)$
 - Algorithm is work-optimal
4. Span
 - Number of levels
 - Span of processLevel (depth of recursion, cost of splitting, base case)
 - Depth of recursion is $\log(n)$
 - Cost of splitting is $\log(n)$
 - Base case is bounded by a constant because we have a constant cutoff
 - Diameter
 - $D(n) = O(d * r * \log(|V| + |E|))$

Conclusion

1. The parallel BFS algorithm is based on binomial heaps