

# Shared Memory Parallel BFS

## Introduction

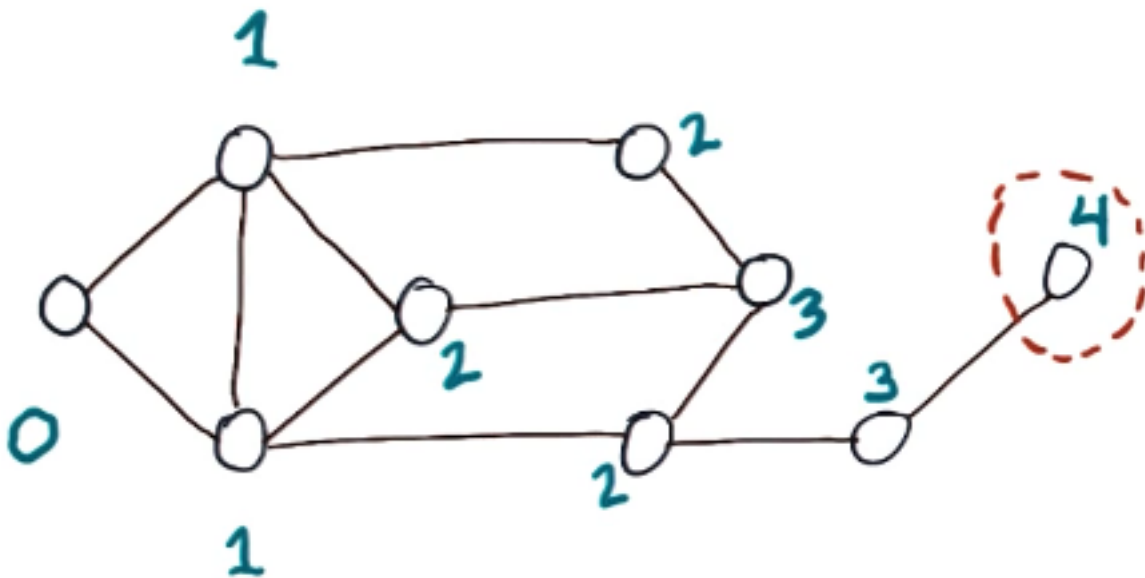
1. Graphs are essential in modern computing
  - Social networks, roads, power grids
2. This lesson examines a parallel breadth-first search algorithm for the dynamic multithreading model

## BFS 101

1. Breadth-first search: Given a graph and a starting vertex, what is the distance to all other nodes in the graph?
2. Work and span
  - $W(n) = O(|V| + |E|)$
  - We visit each vertex at most once and each edge at most once

BFS( $G=(V,E)$ ,  $s$ )

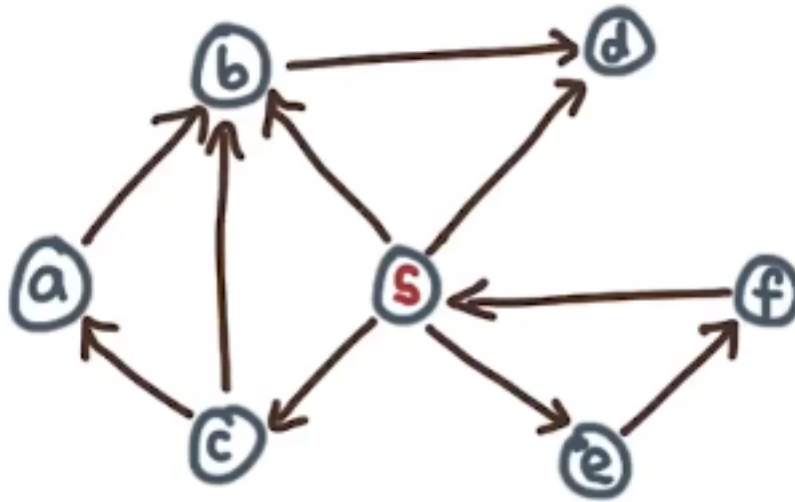
```
{  
  D[V] <- inf  
  D[s] <- 0  
  F <- {s} // queue of unvisited vertices  
  while F != 0 do  
    v <- extract_one(F)  
    for(v,w) <- E do  
      if D[w] = inf then  
        D[w] <- D[v] + 1  
        F <- F U {w}  
  return D // D[x] = dist(s -> x)  
}
```



BFS Example

## BFS Example

1. Consider the following directed graph:



BFS Quiz

---

2. What does F contain right after the third execution of the while loop? There are multiple correct answers due to the ambiguity in “extract\_one”
  - c d e

## Is BFS Inherently Sequential?

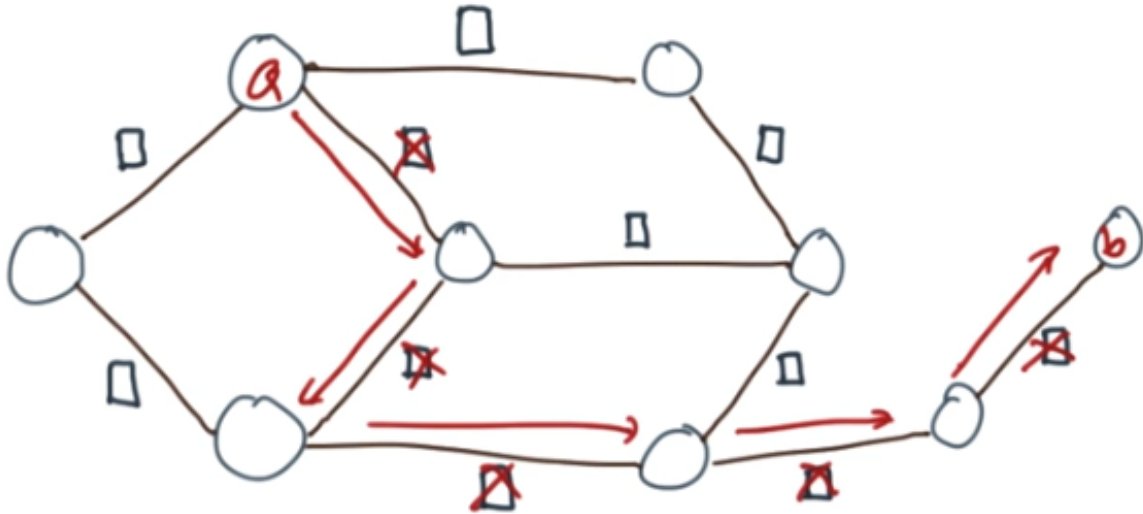
1. BFS Work and Span
  - $W(n) = O(|V| + |E|)$
  - $D(n) = O(|V|)$
  - $W/D = O(1 + |E|/|V|)$
  - In real life, graphs are sparse which means  $|E| = O(|V|)$ 
    - This means the average available parallelism will be a constant, which is no good

## Intuition - Why We Might Do Better

1. The upper bound on the span of BFS shouldn't be the number of vertices, it should be the number of waves
  - Waves are sets of nodes that are reachable from a different set of nodes
  - Diameter: Maximum shortest distance between any pair of vertices in a graph
2. Level-synchronous traversal: Visit nodes level by level
  - Doesn't matter which order we visit nodes within a level
  - This means we can look at nodes simultaneously -> parallelism

## BFS Example 2

1. Find the path with 5 edges in the previous example



BFS Quiz

## High Level Approach to Parallel BFS

1. Two key ideas:
  - Level synchronous
  - Process an entire level in parallel
2. processLevel takes the graph and current frontier and creates a new frontier, as well as updates the distances
  - Frontiers are level-specific waves of the graph
  - l is a level counter
  - Use a special data structure called a bag for the frontiers

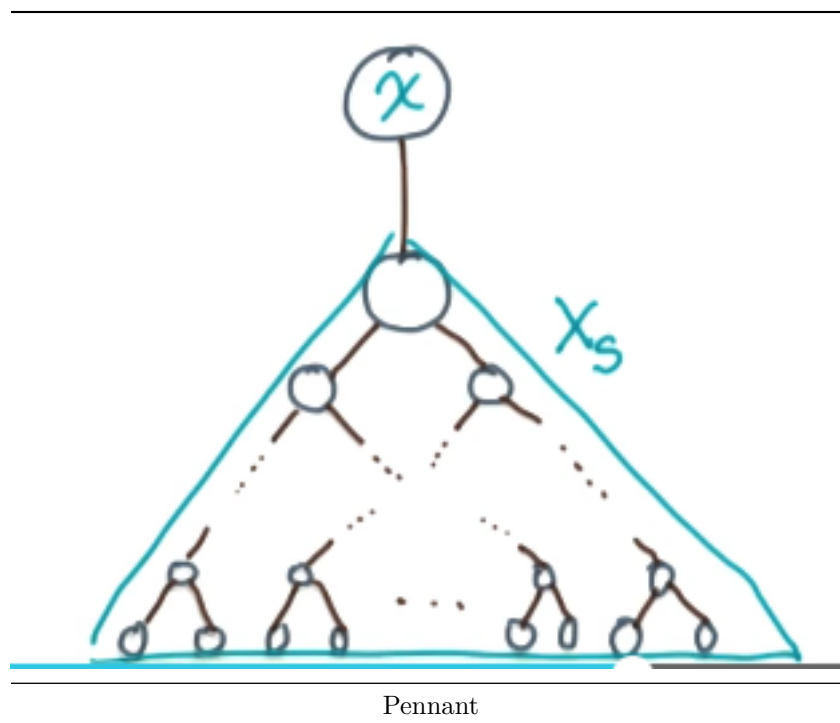
```
parallelBFS(G=(V,E), s)
{
  D[v] <- inf
  D[s] <- 0
  l <- 0
  F0 <- {s}
  while F1 != 0 do
    F(l+1) <- {}
    processLevel(G,F1,F(l+1),D)
    l <- l + 1
  return D
}
```

## Bags: Key Properties

1. Container having the following properties:
  - Unordered collection
  - With repetition
2. Operations
  - Fast traversal
  - Fast and logically associative ( $A \cup B == B \cup A$ )
    - “Union” and “split” should allow us to divide the data structure into roughly equal pieces and combine them back together

## Pennants - Building Blocks for Bags

1. Pennant: Tree with  $2^k$  nodes and a unary root have a child that is a complete binary tree
  - Root:  $x$
  - Child:  $Xs$
2. Consider two same-sized pennants  $x$  and  $y$  with children  $Xs$  and  $Ys$  respectively
  - To combine, make  $x$  the new root,  $y$  the child of  $x$ , and  $Xs$  and  $Ys$  the children of  $y$ 
    - Gives a pennant of size  $2^{(k+1)}$
  - We can undo this operation just as easily by repeating the same steps in reverse



## Pennants

1. What is the output of combining two pennants of different sizes?
  - This is invalid, pennants must be the same size to combine

## Combining Pennants Into Bags

1. Represent 23 in binary = 10111
  - You can use this to split a bag into pennants by creating pennants of size 1, 2, 4, and 16 (in the case of 23)
  - Make an array of pointers to the start of each pennant (called a “spine”)

## Duality Between Bags and Binary Math

1.  $23 + 1 = 10111 + 1 = 11000$
2. If we want to add an element to a bag, we follow the same approach as base-2 addition
  - Carry pennants that are the wrong size into the next slot of the spine

## What is the Cost of Insertion?

1. What is the cost of inserting one element into a bag?

- $O(1)$
  - $O(\log(n))$  (this one)
  - $O(n)$
  - $O(n * \log(n))$
2. In the worst case, we must traverse the entire spine to insert
    - The spine requires  $\text{ceil}(\log_2(n))$  elements to represent an integer  $n$ , so the complexity is  $O(\log(n))$
    - Each insertion requires constant time to perform

## Combining Two Bags

1. What is the cost of bag union?
  - $O(1)$
  - $O(\log(n))$  (this one)
  - $O(n)$
  - $O(n * \log(n))$
  - $O(n^2)$
2. The same logic applies as above
  - This is cool because it takes the same amount of time to insert  $n$  elements into a bag as it does to insert one element
    - The amortized cost is a constant

## Cost of Other Operations

1. What is the cost of bag splitting?
  - $O(1)$
  - $O(\log(n))$  (this one)
  - $O(n)$
  - $O(n * \log(n))$

## Bag Splitting

1. Performing a right bit shift is roughly equivalent to dividing by two
  - $23 = 10111 \gg 1 = 1011 = 11$
2. Cost =  $[O(\log(n) \text{ splits}) * [O(1) \text{ per split}] = O(\log(n))$

## Finishing the Parallel BFS with Bags

1. The following pseudocode implements the processLevel function for the parallel BFS algorithm stated above:

```
processLevel(G, Fl, Fl+1, D)
{
    if |Fl| > phi then
        (A,B) <- bagSplit(Fl)
        spawn processLevel(G, A, Fl+1, D)
        processLevel(G, B, Fl+1, D)
        sync
    else
        for v in Fl do
            parfor (v,w) in E do
                if D[w] = inf then
                    D[w] <- 1 + 1
                    bagInsert(Fl+1, w)
}
```

2. Even though there's a data race in the base case where multiple threads might write to the same  $D[w]$  location, because the algorithm is level-synchronous, each thread is guaranteed to write the same value
3. Work
  - $W(n) = O(|V| + |E|)$ 
    - Algorithm is work-optimal
4. Span
  - Number of levels
  - Span of processLevel (depth of recursion, cost of splitting, base case)
    - Depth of recursion is  $\log(n)$
    - Cost of splitting is  $\log(n)$
    - Base case is bounded by a constant because we have a constant cutoff
  - Diameter
  - $D(n) = O(d * r * \log(|V| + |E|))$

## Conclusion

1. The parallel BFS algorithm is based on binomial heaps